

Pipelines

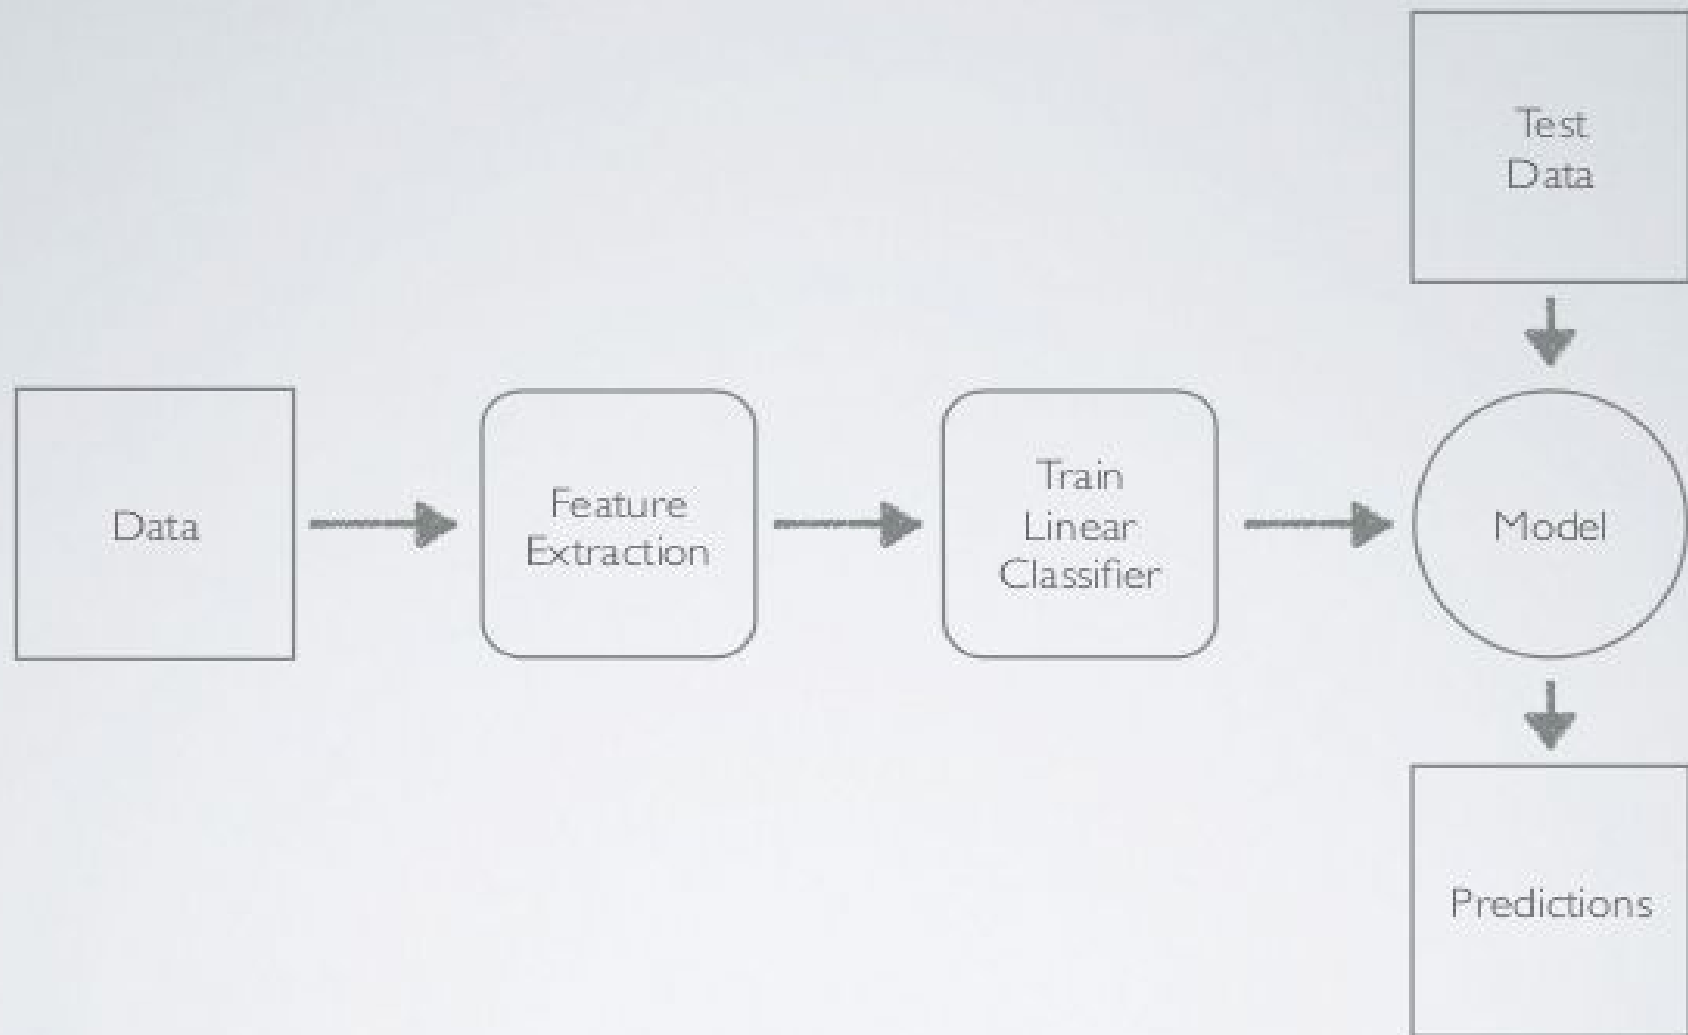
Justin Hooker

Sy Ahmed

Machine Learning Pipelines

- There are standard workflows in a machine learning project that can be automated
- Utility that allows a linear sequence of data transformations to be chained together
- Great way to organize models





Implementing Multinomial NB on text data

```
train = read_file('data/train.tsv')
train_y = extract_targets(train)
train_essays = extract_essays(train)
train_tokens = get_tokens(train_essays)
train_features = extract_features(train)
classifier = MultinomialNB()

scores = []
for train_idx, cv_idx in KFold():
    classifier.fit(train_features[train_idx], train_y[train_idx])
    scores.append(model.score(train_features[cv_idx], train_y[cv_idx]))


print("Score: {}".format(np.mean(scores)))
```

Implementing Multinomial NB on text data with scikit-learn Pipeline

```
pipeline = Pipeline([
    ('extract_essays', EssayExtractor()),
    ('counts', CountVectorizer()),
    ('tf_idf', TfidfTransformer()),
    ('classifier', MultinomialNB())
])

train = read_file('data/train.tsv')
train_y = extract_targets(train)
scores = []
train_idx, cv_idx in KFold():
    model.fit(train[train_idx], train_y[train_idx])
    scores.append(model.score(train[cv_idx], train_y[cv_idx]))

print("Score: {}".format(np.mean(scores)))
```



Custom transformers

Transformer - data in data out black box

Stateless Transformer

```
class HourOfDayTransformer(TransformerMixin):  
  
    def transform(self, X, **transform_params):  
        hours = DataFrame(X['datetime'].apply(lambda x: x.hour))  
        return hours  
  
    def fit(self, X, y=None, **fit_params):  
        return self
```

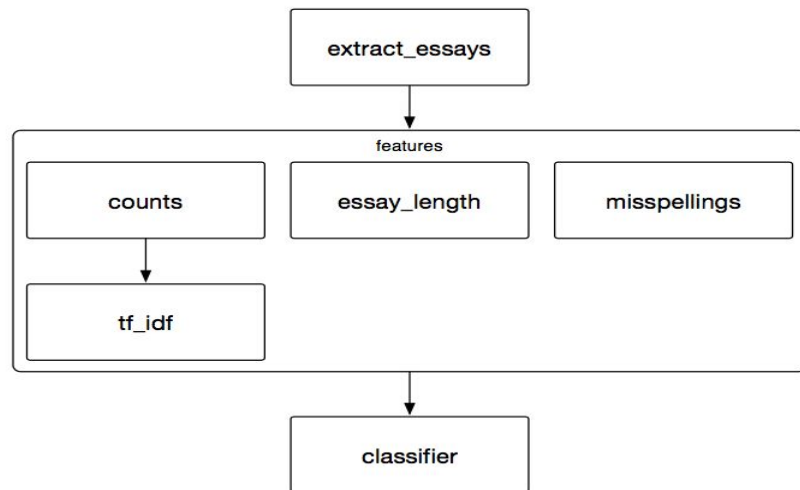
Data Leakage

Leaking knowledge of of the whole dataset to the algorithm

Pipelines help you prevent data leakage in your test harness by ensuring that data preparation like standardization is constrained to each fold of your cross validation procedure.

Feature Unions

```
pipeline = Pipeline([
    ('extract_essays', EssayExtractor()),
    ('features', FeatureUnion([
        ('ngram_tf_idf', Pipeline([
            ('counts', CountVectorizer()),
            ('tf_idf', TfidfTransformer())
        ])),
        ('essay_length', LengthTransformer()),
        ('misspellings', MisspellingCountTransformer())
    ])),
    ('classifier', MultinomialNB())
])
```



Final Thoughts

- The last step of a pipeline is assumed to be the final estimator, so predict method is called instead of transform. For this reason any post processing is done outside of the pipeline.
- Pipelines do not support fit_partial API, so out-of-core training is not possible.
- 2 Important use cases:
 - Data preparation and modeling
 - Feature extraction and feature union
- Adapting, improving models and hyperparameter gridsearch

Demo

Resources

<http://zacstewart.com/2014/08/05/pipelines-of-featureunions-of-pipelines.html>

<https://machinelearningmastery.com/automate-machine-learning-workflows-pipelines-python-scikit-learn/>